

LoRA rank sweep (full data; r = 4–64)

Training data fraction (r = 16)

